

# Machine Learning and Statistical Learning Lecture 1

Daniel Yu

September 9, 2024

---

## Contents

<b>1</b>	<b>Review of Machine Learning Basics</b>	<b>3</b>
1.1	Supervised Learning . . . . .	4
1.2	Regression Model Assumptions . . . . .	4
1.3	Expected Prediction Error . . . . .	4
1.3.1	Squared Error Loss . . . . .	5
1.3.2	Absolute Error . . . . .	5
1.4	Classification . . . . .	6
<b>2</b>	<b>Bias-Variance Tradeoff</b>	<b>6</b>
2.1	Mathematical Decomposition of Expected Prediction Error at fixed $\vec{x}^{(0)}$ . . . . .	7
<b>3</b>	<b>KNN</b>	<b>8</b>
<b>4</b>	<b>Linear Regression - OLS</b>	<b>8</b>
<b>5</b>	<b>Ridge Regression</b>	<b>8</b>
<b>6</b>	<b>Readings</b>	<b>8</b>

---

# 1 Review of Machine Learning Basics

## Statistical Learning

- Input (aka feature/predictor):  $\vec{x} = (x_1, \dots, x_d) = \begin{pmatrix} x_1 \\ \vdots \\ x_d \end{pmatrix} \in X = \mathbb{R}^d$
- Output (label):  $y \in C \subseteq \mathbb{R}$

Note: here  $y$  is an outcome we wish to predict from  $\vec{x}$ .

**Remark.** Measuring a pair  $(\vec{x}, y)$  is a sample of pair of random variables  $R.V.$   $(\vec{X}, Y)$  with underlying distributions.

We want to learn relation between  $\vec{x}$  and  $y$  from training set  $D = \{(\vec{x}^i, y^i), i = 1, \dots, n\}$

Statistical Learning  $\approx$  Machine Learning only that Statistical learning is less focused on algorithms than ML.

**Definition 1.** A learning algorithm is a function that takes training set  $D = \{(\vec{x}^i, y^i), i = 1, \dots, n\}$  as input and has as output a prediction  $y = \hat{f}(\vec{x})$  that for every  $\vec{x}$  predicts  $y$  value.

**Definition 2.** A probability model is a joint probability distribution  $P$  of  $\vec{x} \in X$  and  $y \in C$  on the pair  $(\vec{X}, Y) \in X \times C = \mathbb{R}^d \times C$  of a random vector  $\vec{X} \subseteq \mathbb{R}^d$  and random variable  $Y \in C$ .

**Definition 3.** A function model is a single function  $f : X \rightarrow C$  or a class  $F$  of such functions where one function is assumed to give a good prediction  $y \in C$  from  $\vec{x} \in X$ .

1. Supervised Learning (regression and classification)
2. Unsupervised Learning

**Remark.** One example of Supervised Learning is in generative models:  $(y, \vec{x}) \rightarrow P(y|\vec{x})P(\vec{x})$  note that this is just the chain rule in probability where  $P(x, y) = P(y|x) P(x)$  where  $P(y|\vec{x})$  is the conditional probability that a set of inputs  $\vec{x}$  produces  $y$  and  $P(\vec{x})$  represents the marginal probability of the input features  $\vec{x}$  are in a dataset (i.e. how likely different feature combinations are as each  $x_i$  is a feature).

However, in discriminant models, the above does not apply as they focus solely on learning  $P(y|\vec{x})$  the conditional probability.

One example of Unsupervised Learning is probabilistic modeling:  $(\vec{x}) \rightarrow P(\vec{x})$ , so we are simply estimating the probability distribution of underlying data, using

---

each datapoint  $\vec{x}$  to improve our estimation of the probability distribution  $P(\vec{x})$ .

## 1.1 Supervised Learning

1. regression -  $y \in \mathbb{R}$  is continuous, quantitative
2. classification -  $y \in 1, \dots, K$  is discrete, qualitative

In a statistical learning (theoretical standpoint) they are the same but in a ML algorithmic standpoint they are different!

**Remark.** A probability model has fixed joint prob distribution:  $(\vec{X}, Y) \sim P(\vec{x}, y)$ . For some measurable set  $A \subset \mathbb{R} \times C$ , the probability that  $(\vec{X}, Y) \in A$  is:

$$P(A) = P((\vec{X}, Y) \in A).$$

TODO: Missing stuff – parametric vs non-parametric

## 1.2 Regression Model Assumptions

$$y^{(i)} = h(\vec{x}^{(i)}) + \varepsilon_i.$$

Here  $h(\vec{x})$  depends on some finite or infinite collection of parameters  $\vec{\theta}$

*We assume that errors are random variables of the form  $\varepsilon_i \sim \text{Normal}(0, \sigma^2)$ .*

**Remark.** This is a strong assumption!

*Proof.*

$$E(Y|\vec{X} = \vec{x}) = E(h(\vec{X}) + \varepsilon|\vec{X} = \vec{x}) \text{ since } h(\vec{X}) \text{ and } \varepsilon \text{ are independent} \quad (1)$$

$$= E(h(\vec{X})|\vec{X} = \vec{x}) + E(\varepsilon|\vec{X} = \vec{x}) \quad (2)$$

$$= h(\vec{x}) + 0 \quad (3)$$

$$= h(\vec{x}) \quad (4)$$

□

**Remark.** Proof above is only true for the assumption that  $\varepsilon$  is random white noise.

## 1.3 Expected Prediction Error

Assume we have an algorithm for estimating  $h(\vec{x}^0)$  for test point  $\vec{x}^0$ . The estimator is:

$$\bar{y}^0 = \hat{h}(\vec{x}^0) \approx h(\vec{x}^0).$$

**Definition 4.** Our loss function is then defined as:

$$L(Y, h(\vec{X})).$$

**Loss Function Types**

- $L(Y, h(\vec{X})) = (Y - h(\vec{X}))^2$  square loss error
- $L(Y, h(\vec{X})) = |Y - h(\vec{X})|$  absolute loss error

**Definition 5.** Expected Prediction Error (Expected Value of Error)

$$EPE(h) = E_{Y, \vec{X}}[L(Y, h(\vec{X}))] = E_{\vec{X}} E_{Y|\vec{X}}[L(Y, h(\vec{X})|\vec{X})].$$

We want to minimize pointwise:

$$\hat{h}(\vec{x}) = \operatorname{argmin}_c E_{Y|\vec{X}}[L(Y, c)|\vec{X} \rightarrow \vec{x}].$$

where  $c = h(\vec{X})$

**Remark.** In Theory we want prediction error for all future values but in practice we only have to the test error.

### 1.3.1 Squared Error Loss

Let  $L(Y, h(\vec{X})) = (Y - h(\vec{X}))^2$  square loss error

$$\begin{aligned} \hat{h}(\vec{x}) &= \operatorname{argmin}_c E_{Y|\vec{X}}[(Y - c)^2|\vec{X} = \vec{x}]. \\ &= \operatorname{argmin}_c E_{Y|\vec{X}}[Y^2 - 2cY + c^2|\vec{X} = \vec{x}]. \end{aligned}$$

Note:

$$\begin{aligned} E[Y^2 - 2cY + c^2|\vec{X} = \vec{x}] &= E[Y]^2 - 2cE[Y] + c^2. \\ \frac{d}{dc}[E[Y]^2 - 2cE[Y] + c^2] &= -2E[Y] + 2c = 0. \\ c &= E[Y] = E[Y|\vec{X} = \vec{x}]. \end{aligned}$$

Logically this makes sense because the minimizer of the squared loss function when  $\hat{h}(\vec{x}) = E[Y|\vec{X} = \vec{x}]$  is the mean (expected value) of Y given X = x which captures the squared loss function's central tendency as it squares distance.

### 1.3.2 Absolute Error

Let  $L(Y, h(\vec{X})) = |Y - h(\vec{X})|$  absolute loss error:

$$\hat{h}(\vec{x}) = \operatorname{argmin}_c E_{Y|\vec{X}}[|Y - c||\vec{X} = \vec{x}].$$

Expand the  $|Y - c|$  piecewise:

$$E[|Y - c|] = \int_{-\infty}^c (c - y)p(y)dy + \int_c^{\infty} (y - c)p(y)dy.$$

Take the derivative to minimize the expected value:

$$\frac{d}{dc}(E[|Y - c|]) = \frac{d}{dc}[\int_{-\infty}^c (c - y)p(y)dy + \int_c^{\infty} (y - c)p(y)dy] = 0.$$

---


$$= \int_{-\infty}^c p(y)dy - \int_c^{\infty} p(y)dy = 0.$$

$$\int_{-\infty}^c p(y)dy = \int_c^{\infty} p(y)dy.$$

This implies that to minimize the Expected Value,  $c$  is the median of the distribution of  $Y$ , because the cumulative probability mass to the left of  $c$  must be equal to the cumulative probability mass to the right of  $c$ .

$$\hat{h}(\vec{x}) = \text{median}(Y|\vec{X} = \vec{x}).$$

**Remark.** Absolute Value is not differentiable (although it is continuous) so greedy descent methods can't be used.

## 1.4 Classification

Consider the binary case where  $Y = 0, 1$  for class 1 and class 2 respectively:

$$\begin{aligned}\hat{h}(\vec{x}) &= E[Y|\vec{X} = \vec{x}] \\ &= P(\text{class 1}|\vec{X} = \vec{x}) \cdot 1 + P(\text{class 2}|\vec{X} = \vec{x}) \cdot 0 \\ &= P(Y = 1|\vec{X} = \vec{x}) \cdot 1 + P(Y = 0|\vec{X} = \vec{x}) \cdot 0 \\ &= P(Y = 1|\vec{X} = \vec{x}) \in \mathbb{R}.\end{aligned}$$

In the multiclass case:  $Y = 1, \dots, K$ :

$$\begin{aligned}\hat{x} &= \arg \min_{f(x)} E_{Y|\vec{X}}[L(Y, h(X))|\vec{X} = \vec{x}] \\ &= \arg \min_k E_{Y|\vec{X}}[L(Y, k)|\vec{X} = \vec{x}] \\ &= \arg \min_k \sum_{y=1}^K L(y, k)P(Y = y|\vec{X} = \vec{x}).\end{aligned}$$

TODO classifiers stuff

## 2 Bias-Variance Tradeoff

Assuming, iid  $\varepsilon_i$  as noise:

**Definition 6.** The Expected Prediction (in practice test) Error (EPE) for a fixed test point  $\vec{x}^0$  is the average over both  $y^0$  and the entire training set  $D$ .

$$EPE(\vec{x}^0) = E_{y^0, D}[L(y^0, \hat{h}(\vec{x}^0))].$$

We want in theory **full expected prediction error**, meaning we average over all training sets:  $D = \{(\vec{x}^i, \vec{y}^i), i = 1, \dots, n\}$  where each  $D_i$  is one possible split of train/test data and over all possible values  $Y = y^0$  at  $\vec{x}^0$ .

---

**Remark.**  $D$  is a random variable whose distribution is all possible train/test splits.  $y^0$  is a random variable representing the true value of the outcome at  $\vec{x}^0$  but inherits the random "noise" in the data generation process, leading to different  $\hat{y}^0$  being observed.

## 2.1 Mathematical Decomposition of Expected Prediction Error at fixed $\vec{x}^{(0)}$

**Note.** This is for the case where the loss function is the **Mean Squared Error** loss function.

$$\begin{aligned} EPE(\vec{x}^0) &= E_{y^0, D}[(y^0 - \hat{y}^0)^2] \\ &= E_{y^0} E_D[(y^0 - \hat{y}^0)^2] \end{aligned}$$

$$\begin{aligned} &= E_{y^0} E_D\{[(y^0 - E_{y^0}[y^0]) + (E_{y^0}[y^0] - E_D[\hat{y}^0]) + (E_D[\hat{y}^0] - \hat{y}^0)]^2\}. \\ &= E_{y^0}[y^0 - E_{y^0}[y^0]]^2 + [E_{y^0}[y^0] - E_D[\hat{y}^0]]^2 + E_D[E_D[\hat{y}^0] - \hat{y}^0]^2. \end{aligned}$$

where  $y_0 = h(\vec{x}^0) + \varepsilon$  and  $\hat{y}^0 = \hat{h}(\vec{x})$

**Note.** Note all the cross terms dissappear from the expression about which is why we can eliminate them:

$$\begin{aligned} &E_{y^0} E_D[(Y^0 - E_{y^0}[y^0]) (E_D[\hat{y}^0] - \hat{y}^0)] \text{ the first expression does not depend on } D \\ &= E_{y^0}[(y^0 - E_{y^0}[y^0]) E_D(E_D[\hat{y}^0] - \hat{y}^0)] \\ &= 0. \end{aligned}$$

The last step is possible because  $y^0, \hat{y}^0$  are independent since  $\varepsilon_i$  are independent and  $\hat{y}^0$  depends on  $D$ , and by probability  $E[X \cdot Y] = E[X] \cdot E[Y]$  when  $X, Y$  are independent.

1.  $(y^0 - E_{y^0}[y^0])$  – can be thought of as the **noise** of the model that distorts the true value of  $y$  from its expected value (coming from the  $\varepsilon$  assumed to be random gaussian noise)
2.  $E_{y^0}[y^0] - E_D[\hat{y}^0]$  – can be thought of as the **bias**, difference between true expected output  $E_{y^0}[y^0]$  and the model's prediction  $E_D[\hat{y}^0]$ .
3.  $(E_D[\hat{y}^0] - \hat{y}^0)$  – can be thought of as **variance**, variability of model's prediction  $\hat{y}^0$  from the actual expected value  $E_D[\hat{y}^0]$  (recall that  $D$  is the set of all possible training sets from different train/test splits where our model takes only train set).

We can rewrite the formula since  $var = \frac{1}{N} \sum_{i=1}^N (x_i - \mu)^2$ :

$$\begin{aligned} &= var_{y^0}(y^0) + [E_{y^0}[y^0] - E_D[\hat{y}^0]]^2 + var_D(\hat{y}^0) \\ &= \sigma^2 + [h(x^0) - E_D[\hat{y}^0]]^2 + var_D(\hat{y}^0) \quad \text{from } N(\mu = h(x^0), \sigma) \\ &= \sigma^2 + [bias]^2 (\hat{y}^0) + var_D(\hat{y}^0) \\ &= \text{unavoidable error} + \text{bias}^2 + var_D(\hat{y}^0) \end{aligned}$$

---

where bias is difference between true and predicted  
TODO knn and LINEAR model notes

### 3 KNN

### 4 Linear Regression - OLS

### 5 Ridge Regression

Motivation: no longer working with unbiased estimators, trade off some bias to large decrease in variance.

**Definition 7.** Assume data  $D = (X, \vec{y})$  is centered, the mean  $E(X) = \vec{0}$  and  $E(\vec{y}) = 0$ , so

$$h(\vec{x}) = \vec{\theta}^T \vec{x} = \theta_1 x_1 + \dots + \theta_d x_d.$$

**Ridge Regression** minimize cost function:

$$J(\vec{\theta}) = \sum_{i=1}^n (y^i - h_{\theta}(\vec{x}^i))^2 + \lambda \sum_{j=1}^d \theta_j^2 = (X\vec{\theta} - \vec{y})^T (X\vec{\theta} - \vec{y}) + \lambda \vec{\theta}^T \vec{\theta}.$$

Taking the first derivative:

$$\begin{aligned} \frac{d}{d\vec{\theta}} J &= \dots \\ &= \dots \end{aligned}$$

$$\vec{\theta} = (X^T X + \lambda I)^{-1} X^T \vec{y}.$$

TODO More examples: TODO High Dimension Data and Reduction

### 6 Readings

1. <https://mlu-explain.github.io/bias-variance/>